


# Kafka Producer and Consumer Examples Using Java

by Gaurav Garg · May. 25, 18 · Big Data Zone · Tutorial



Introducing New Cisco Certified DevNet Specialist: IoT  
 Validate your IoT application development knowledge with the practice of Cisco IoT edge compute and architecture. ▶

In my last article, we discussed how to setup Kafka using Zookeeper. In this article, we will see how to produce and consume records/messages with Kafka brokers.

Before starting with an example, let's get familiar first with the common terms and some commands used in Kafka.

**Record:** Producer sends messages to Kafka in the form of records. A record is a key-value pair. It contains the topic name and partition number to be sent. Kafka broker keeps records inside topic partitions. Records sequence is maintained at the partition level. You can define the logic on which basis partition will be determined.

**Topic:** Producer writes a record on a topic and the consumer listens to it. A topic can have many partitions but must have at least one.

**Partition:** A topic partition is a unit of parallelism in Kafka, i.e. two consumers cannot consume messages from the same partition at the same time. A consumer can consume from multiple partitions at the same time.

**Offset:** A record in a partition has an offset associated with it. Think of it like this: partition is like an array; offsets are like indexes.

**Producer:** Creates a record and publishes it to the broker.

**Consumer:** Consumes records from the broker.

**Commands:** In Kafka, a setup directory inside the bin folder is a script (kafka-topics.sh), using which, we can create and delete topics and check the list of topics. Go to the Kafka home directory.

- Execute this command to see the list of all topics.
  - `./bin/kafka-topics.sh --list --zookeeper localhost:2181` .

- *localhost:2181* is the Zookeeper address that we defined in the `server.properties` file in the previous article.
- Execute this command to create a topic.
  - **`./bin/kafka-topics.sh --create --zookeeper localhost:2181 --replication-factor 1 --partitions 100 --topic demo .`**
  - `replication-factor` : if Kafka is running in a cluster, this determines on how many brokers a partition will be replicated. The `partitions` argument defines how many partitions are in a topic.
  - After a topic is created you can increase the partition count but it cannot be decreased. `demo` , here, is the topic name.
- Execute this command to delete a topic.
  - **`./bin/kafka-topics.sh --zookeeper localhost:2181 --delete --topic demo .`**
  - This command will have no effect if in the Kafka `server.properties` file, if `delete.topic.enable` is not set to be true.
- Execute this command to see the information about a topic.
  - **`./bin/kafka-topics.sh --describe --topic demo --zookeeper localhost:2181 .`**

Now that we know the common terms used in Kafka and the basic commands to see information about a topic ,let's start with a working example.

```

1 public interface IKafkaConstants {
2     public static String KAFKA_BROKERS = "localhost:9092";
3
4     public static Integer MESSAGE_COUNT=1000;
5
6     public static String CLIENT_ID="client1";
7
8     public static String TOPIC_NAME="demo";
9
10    public static String GROUP_ID_CONFIG="consumerGroup1";
11
12    public static Integer MAX_NO_MESSAGE_FOUND_COUNT=100;
13
14    public static String OFFSET_RESET_LATEST="latest";
15
16    public static String OFFSET_RESET_EARLIER="earliest";
17
18    public static Integer MAX_POLL_RECORDS=1;
19 }

```

The above snippet contains some constants that we will be using further.

```
1 import java.util.Properties;
2
3 import org.apache.kafka.clients.producer.KafkaProducer;
4 import org.apache.kafka.clients.producer.Producer;
5 import org.apache.kafka.clients.producer.ProducerConfig;
6 import org.apache.kafka.common.serialization.LongSerializer;
7 import org.apache.kafka.common.serialization.StringSerializer;
8
9 import com.gaurav.kafka.constants.IKafkaConstants;
0
1 public class ProducerCreator {
2
3     public static Producer<Long, String> createProducer() {
4         Properties props = new Properties();
5         props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, IKafkaConstants.KAFKA_BROKERS);
6         props.put(ProducerConfig.CLIENT_ID_CONFIG, IKafkaConstants.CLIENT_ID);
7         props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, LongSerializer.class.getName());
8         props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
9         //props.put(ProducerConfig.PARTITIONER_CLASS_CONFIG, CustomPartitioner.class.getName());
0         return new KafkaProducer<>(props);
1     }
2 }
```

The above snippet creates a Kafka producer with some properties.

- **BOOTSTRAP\_SERVERS\_CONFIG**: The Kafka broker's address. If Kafka is running in a cluster then you can provide comma (,) separated addresses. For example: localhost:9091,localhost:9092
- **CLIENT\_ID\_CONFIG**: Id of the producer so that the broker can determine the source of the request.
- **KEY\_SERIALIZER\_CLASS\_CONFIG**: The class that will be used to serialize the key object. In our example, our key is `Long`, so we can use the `LongSerializer` class to serialize the key. If in your use case you are using some other object as the key then you can create your custom serializer class by implementing the **Serializer** interface of Kafka and overriding the `serialize` method.
- **VALUE\_SERIALIZER\_CLASS\_CONFIG**: The class that will be used to serialize the value object. In our example, our value is `String`, so we can use the `StringSerializer` class to serialize the key. If your value is some other object then you create your custom serializer class. For example:

```
1 import java.util.Map;
2 import org.apache.kafka.common.serialization.Serializer;
3 import org.slf4j.Logger;
4 import org.slf4j.LoggerFactory;
5
```

```

6 import com.fasterxml.jackson.databind.ObjectMapper;
7 import com.gaurav.kafka.pojo.CustomObject;
8
9 public class CustomSerializer implements Serializer<CustomObject> {
10
11     @Override
12     public void configure(Map<String, ?> configs, boolean isKey) {
13
14     }
15
16     @Override
17     public byte[] serialize(String topic, CustomObject data) {
18         byte[] retVal = null;
19         ObjectMapper objectMapper = new ObjectMapper();
20         try {
21             retVal = objectMapper.writeValueAsString(data).getBytes();
22         } catch (Exception exception) {
23             System.out.println("Error in serializing object"+ data);
24         }
25         return retVal;
26     }
27
28     @Override
29     public void close() {
30
31     }
32 }

```

- **PARTITIONER\_CLASS\_CONFIG**: The class that will be used to determine the partition in which the record will go. In the demo topic, there is only one partition, so I have commented this property. You can create your custom partitioner by implementing the **CustomPartitioner** interface. For example:

```

1 import java.util.Map;
2
3 import org.apache.kafka.clients.producer.Partitioner;
4 import org.apache.kafka.common.Cluster;
5
6 public class CustomPartitioner implements Partitioner{
7
8     private static final int PARTITION_COUNT=50;
9
10    @Override
11    public void configure(Map<String, ?> configs) {
12
13    }
14
15    @Override
16    public int partition(String topic, Object key, byte[] keyBytes, Object value, byte[] valueBytes, Cluster clu
17        Integer keyInt=Integer.parseInt(key.toString());
18        return keyInt % PARTITION_COUNT;
19    }
20
21    @Override
22    public void close() {
23
24    }
25

```

```
6}
```

In above the `CustomPartitioner` class, I have overridden the method **partition** which returns the partition number in which the record will go.

```
1 import java.util.Collections;
2 import java.util.Properties;
3
4 import org.apache.kafka.clients.consumer.Consumer;
5 import org.apache.kafka.clients.consumer.ConsumerConfig;
6 import org.apache.kafka.clients.consumer.KafkaConsumer;
7 import org.apache.kafka.common.serialization.LongDeserializer;
8 import org.apache.kafka.common.serialization.StringDeserializer;
9
10 import com.gaurav.kafka.constants.IKafkaConstants;
11
12 public class ConsumerCreator {
13
14     public static Consumer<Long, String> createConsumer() {
15         Properties props = new Properties();
16         props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, IKafkaConstants.KAFKA_BROKERS);
17         props.put(ConsumerConfig.GROUP_ID_CONFIG, IKafkaConstants.GROUP_ID_CONFIG);
18         props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG, LongDeserializer.class.getName());
19         props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG, StringDeserializer.class.getName());
20         props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, IKafkaConstants.MAX_POLL_RECORDS);
21         props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, "false");
22         props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, IKafkaConstants.OFFSET_RESET_EARLIER);
23
24         Consumer<Long, String> consumer = new KafkaConsumer<>(props);
25         consumer.subscribe(Collections.singletonList(IKafkaConstants.TOPIC_NAME));
26         return consumer;
27     }
28
29 }
```

The above snippet creates a Kafka consumer with some properties.

- **BOOTSTRAP\_SERVERS\_CONFIG**: The Kafka broker's address. If Kafka is running in a cluster then you can provide comma (,) seperated addresses. For example: `localhost:9091,localhost:9092`.
- **GROUP\_ID\_CONFIG**: The consumer group id used to identify to which group this consumer belongs.
- **KEY\_DESERIALIZER\_CLASS\_CONFIG**: The class name to deserialize the key object. We have used `Long` as the key so we will be using **LongDeserializer** as the deserializer class. You can create your custom deserializer by implementing the **Deserializer** interface provided by Kafka.
- **VALUE\_DESERIALIZER\_CLASS\_CONFIG**: The class name to deserialize the value object. We have used `String` as the value so we will be using **StringDeserializer** as the deserializer class. You can create your custom deserializer. For example:

```

1 import java.util.Map;
2
3 import org.apache.kafka.common.serialization.Deserializer;
4 import org.slf4j.Logger;
5 import org.slf4j.LoggerFactory;
6
7 import com.fasterxml.jackson.databind.ObjectMapper;
8 import com.gaurav.kafka.pojo.CustomObject;
9
0 public class CustomObjectDeserializer implements Deserializer<CustomObject> {
1
2     @Override
3     public void configure(Map<String, ?> configs, boolean isKey) {
4     }
5
6     @Override
7     public CustomObject deserialize(String topic, byte[] data) {
8         ObjectMapper mapper = new ObjectMapper();
9         CustomObject object = null;
0         try {
1 object = mapper.readValue(data, CustomObject.class);
2         } catch (Exception exception) {
3 System.out.println("Error in deserializing bytes "+ exception);
4         }
5         return object;
6     }
7
8     @Override
9     public void close() {
0     }
1
2 }

```

- **MAX\_POLL\_RECORDS\_CONFIG**: The max count of records that the consumer will fetch in one iteration.
- **ENABLE\_AUTO\_COMMIT\_CONFIG**: When the consumer from a group receives a message it must commit the offset of that record. If this configuration is set to be true then, periodically, offsets will be committed, but, for the production level, this should be false and an offset should be committed manually.
- **AUTO\_OFFSET\_RESET\_CONFIG**: For each consumer group, the last committed offset value is stored. This configuration comes handy if no offset is committed for that group, i.e. it is the new group created.
  - Setting this value to **earliest** will cause the consumer to fetch records from the beginning of offset i.e from zero.
  - Setting this value to **latest** will cause the consumer to fetch records from the new records. By new records mean those created after the consumer group became active.

```
1 import java.util.concurrent.ExecutionException;
2
3 import org.apache.kafka.clients.consumer.Consumer;
4 import org.apache.kafka.clients.consumer.ConsumerRecords;
5 import org.apache.kafka.clients.producer.Producer;
6 import org.apache.kafka.clients.producer.ProducerRecord;
7 import org.apache.kafka.clients.producer.RecordMetadata;
8
9 import com.gaurav.kafka.constants.IKafkaConstants;
0 import com.gaurav.kafka.consumer.ConsumerCreator;
1 import com.gaurav.kafka.producer.ProducerCreator;
2
3 public class App {
4     public static void main(String[] args) {
5         runProducer();
6         //runConsumer();
7     }
8
9     static void runConsumer() {
0         Consumer<Long, String> consumer = ConsumerCreator.createConsumer();
1
2         int noMessageFound = 0;
3
4         while (true) {
5             ConsumerRecords<Long, String> consumerRecords = consumer.poll(1000);
6             // 1000 is the time in milliseconds consumer will wait if no record is found at broker.
7             if (consumerRecords.count() == 0) {
8                 noMessageFound++;
9                 if (noMessageFound > IKafkaConstants.MAX_NO_MESSAGE_FOUND_COUNT)
0                     // If no message found count is reached to threshold exit loop.
1                     break;
2                 else
3                     continue;
4             }
5
6             //print each record.
7             consumerRecords.forEach(record -> {
8                 System.out.println("Record Key " + record.key());
9                 System.out.println("Record value " + record.value());
0                 System.out.println("Record partition " + record.partition());
1                 System.out.println("Record offset " + record.offset());
2             });
3
4             // commits the offset of record to broker.
5             consumer.commitAsync();
6         }
7         consumer.close();
8     }
9
0     static void runProducer() {
1 Producer<Long, String> producer = ProducerCreator.createProducer();
2
3         for (int index = 0; index < IKafkaConstants.MESSAGE_COUNT; index++) {
4             ProducerRecord<Long, String> record = new ProducerRecord<Long, String>(IKafkaConstants.TOPIC_NAME,
5 "This is record " + index);
6             try {
7                 RecordMetadata metadata = producer.send(record).get();
8                 System.out.println("Record sent with key " + index + " to partition " + metadata.parti
9 + " with offset " + metadata.offset());
0             }
1             catch (ExecutionException e) {
2                 System.out.println("Error in sending record");
3                 System.out.println(e);
4             }
5         }
6     }
7 }
8
9
0
1
2
3
```

```

4         }
5         catch (InterruptedException e) {
6             System.out.println("Error in sending record");
7             System.out.println(e);
8         }
9     }
0 }
1}

```

The above snippet explains how to produce and consume messages from a Kafka broker. If you want to run a producer then call the **runProducer** function from the main function. If you want to run a consumer, then call the **runConsumer** function from the main function.

- The offset of records can be committed to the broker in both asynchronous and synchronous ways. Using the synchronous way, the thread will be blocked until an offset has not been written to the broker.

## Conclusion

We have seen how Kafka producers and consumers work. You can check out the whole project on my GitHub page. If you are facing any issues with Kafka, please ask in the comments. In next article, I will be discussing how to set up monitoring tools for Kafka using Burrow.

Accelerate your analytics workflow with OmniSci. [Learn more about our platform](#) that delivers zero-latency querying and visual exploration of your data.

Presented by OmniSci

## Like This Article? Read More From DZone



**A Beginner's Guide to Apache Kafka**



**Using Apache Kafka for Real-Time Event Processing**



**Kafka Producer Delivery Semantics**



**Free DZone Refcard  
Understanding Apache Spark Failures and Bottlenecks**

Topics: [APACHE KAFKA](#) , [BIG DATA](#) , [BIG DATA PLATFOMS](#)

Opinions expressed by DZone contributors are their own.